

## Exercise: Try different MLP architectures on MNIST dataset.

### Step By Step Process

- Lets start working on this Assignments and as we know our main objective in assignment to implement Multi Layer Perceptron using Keras on MNIST Dataset And in this we will work with Multi Layer Perceptron and also MLP with Dropout and Batch Normalization.
- First we will start with importing multiple important libraries which are being used in this assignments and after that we will load our MNIST Dataset.
- After loading MNIST Dataset with 100000 data points and then we will split our dataset into train and test and we know our dataset i.e image is in the shape of (28 X 28) so we will first reshape our dataset i.e we will convert the (2828) *vector into single dimensional vector of 1 784* .
- Now in this step we will do data point normalization and as we observe the our data in which each cell is having a value between 0-255 so before we move to apply machine learning algorithms lets try to normalize the data.
- And as we know our class labels in this dataset is in numbers between {0,1,2,3.....9} and now we will convert it into one hot encoded vector
- After doing all above now we will start working with our models in this we will work with softmax classifier with multiple hidden layers that means we will work with softmax with multiple hidden layers and try to observe the performance by changing the no of layers and in this we will also work with dropout and batch normalization.
- And after doing all this we will try to observe the performance of train and test val so that we will be able to know our model should not overfit.

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal

# importing other lib
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
```

```
C:\Users\nisha\Anaconda3\lib\site-packages\h5py\__init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.
  from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

In [14]:

```
def plt_dynamic_model(x, vy, ty):
    plt.figure(figsize=(10,5))
    plt.plot(x, vy, 'b', label="val Loss")
    plt.plot(x, ty, 'r', label="Train Loss")
    plt.xlabel('Epochs val')
    plt.ylabel('Categorical Crossentropy Loss')
    plt.title('\nCategorical Crossentropy Loss VS Epochs')
    plt.legend()
    plt.show()
```

### Loading MNIST Dataset

- And then splitting it into train and test

In [2]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
print("Number of training examples:",X_train.shape[0],"and each image is of shape (%d, %d)"%(X_train.shape[1],X_train.shape[2]))
print("Number of training examples :",X_test.shape[0],"and each image is of shape (%d,
```

```
%d)"%(X_test.shape[1],X_test.shape[2]))
```

Number of training examples: 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

In [3]:

```
# if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])

# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape
(%d)"%(X_train.shape[1]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape
(%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of test examples : 10000 and each image is of shape (784)

In [4]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
 0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  136 253 253 253 212 135 132  16
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

## Data point normalization

In [5]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - Xmin)/(Xmax-Xmin) = X/255$ 
```

```
X_train = X_train/255
```

$$\bar{X}_{\text{test}} = \bar{X}_{\text{test}}/255$$

```
# example data point after normlizing
```

```
print(X_train[0])
```

[illegible]



- As we know our class labels in this dataset is in numbers between {0,1,2,3.....9} and now we will convert it into one hot encoded vector

In [6]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5  
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

#

In [31]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.initializers import he_normal
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

## 1. Softmax classifier with 2-hidden layers Without dropout and Batch Normalization

In [8]:

```
model2 = Sequential()

model2.add(Dense(385, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model2.add(Dense(120, activation='relu', kernel_initializer=he_normal(seed=None)))
model2.add(Dense(output_dim, activation='softmax'))

# model Summary
print("Model Summary :- \n",model2.summary())

# Compiling the model
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history2 = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type)    | Output Shape | Param # |
|-----------------|--------------|---------|
| dense_4 (Dense) | (None, 385)  | 302225  |
| dense_5 (Dense) | (None, 120)  | 46320   |
| dense_6 (Dense) | (None, 10)   | 1210    |

=====  
 Total params: 349,755  
 Trainable params: 349,755  
 Non-trainable params: 0

---

#### Model Summary :-


None  
Train on 60000 samples, validate on 10000 samples  
Epoch 1/20  
60000/60000 [=====] - 9s 142us/step - loss: 0.2441 - acc: 0.9290 -  
val\_loss: 0.1123 - val\_acc: 0.9653  
Epoch 2/20  
60000/60000 [=====] - 8s 127us/step - loss: 0.0918 - acc: 0.9718 -  
val\_loss: 0.0919 - val\_acc: 0.9717  
Epoch 3/20  
60000/60000 [=====] - 8s 132us/step - loss: 0.0573 - acc: 0.9827 -  
val\_loss: 0.0746 - val\_acc: 0.9765  
Epoch 4/20  
60000/60000 [=====] - 7s 123us/step - loss: 0.0389 - acc: 0.9877 -  
val\_loss: 0.0668 - val\_acc: 0.9798  
Epoch 5/20  
60000/60000 [=====] - 7s 123us/step - loss: 0.0304 - acc: 0.9905 -  
val\_loss: 0.0685 - val\_acc: 0.9799  
Epoch 6/20  
60000/60000 [=====] - 7s 122us/step - loss: 0.0226 - acc: 0.9932 -  
val\_loss: 0.0684 - val\_acc: 0.9810  
Epoch 7/20  
60000/60000 [=====] - 7s 122us/step - loss: 0.0185 - acc: 0.9941 -  
val\_loss: 0.0821 - val\_acc: 0.9773  
Epoch 8/20  
60000/60000 [=====] - 9s 142us/step - loss: 0.0157 - acc: 0.9948 -  
val\_loss: 0.0779 - val\_acc: 0.9794  
Epoch 9/20  
60000/60000 [=====] - 8s 129us/step - loss: 0.0162 - acc: 0.9945 -  
val\_loss: 0.0768 - val\_acc: 0.9794  
Epoch 10/20  
60000/60000 [=====] - 8s 128us/step - loss: 0.0103 - acc: 0.9966 -  
val\_loss: 0.0705 - val\_acc: 0.9818  
Epoch 11/20  
60000/60000 [=====] - 8s 127us/step - loss: 0.0081 - acc: 0.9972 -  
val\_loss: 0.0818 - val\_acc: 0.9794  
Epoch 12/20  
60000/60000 [=====] - 8s 126us/step - loss: 0.0145 - acc: 0.9951 -  
val\_loss: 0.0845 - val\_acc: 0.9788  
Epoch 13/20  
60000/60000 [=====] - 8s 131us/step - loss: 0.0088 - acc: 0.9971 -  
val\_loss: 0.0924 - val\_acc: 0.9790  
Epoch 14/20  
60000/60000 [=====] - 7s 122us/step - loss: 0.0090 - acc: 0.9970 -  
val\_loss: 0.0959 - val\_acc: 0.9797  
Epoch 15/20  
60000/60000 [=====] - 9s 145us/step - loss: 0.0072 - acc: 0.9976 -  
val\_loss: 0.0945 - val\_acc: 0.9804  
Epoch 16/20  
60000/60000 [=====] - 7s 120us/step - loss: 0.0056 - acc: 0.9982 -  
val\_loss: 0.0855 - val\_acc: 0.9816  
Epoch 17/20  
60000/60000 [=====] - 8s 132us/step - loss: 0.0093 - acc: 0.9969 -  
val\_loss: 0.1034 - val\_acc: 0.9781s - 1  
Epoch 18/20  
60000/60000 [=====] - 8s 125us/step - loss: 0.0095 - acc: 0.9967 -  
val\_loss: 0.0933 - val\_acc: 0.9796  
Epoch 19/20  
60000/60000 [=====] - 8s 134us/step - loss: 0.0050 - acc: 0.9984 -  
val\_loss: 0.0809 - val\_acc: 0.9809  
Epoch 20/20  
60000/60000 [=====] - 7s 124us/step - loss: 0.0104 - acc: 0.9967 -  
val\_loss: 0.0834 - val\_acc: 0.9820

In [22]:

```
x = list(range(1,nb_epoch+1))  
  
# getting Val loss  
vy = history2.history['val_loss']  
# getting Train loss  
ty = history2.history['loss']  
  
# function call  
plt_dynamic_model(x, vy, ty)
```

```
# Evaluating the model
model_score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', model_score[0])
print('Test accuracy:', model_score[1])

# saving train and test accuracy of the model
model2_test_acc = model_score[1]
model2_train_acc = history2.history['acc']
```



Test score: 0.08338531920399564  
Test accuracy: 0.982

## Softmax classifier with 2-hidden layers With dropout and Batch Normalization

In [9]:

```
model2dn = Sequential()

model2dn.add(Dense(375, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model2dn.add(BatchNormalization())
model2dn.add(Dropout(0.5))

model2dn.add(Dense(112, activation='relu', kernel_initializer=he_normal(seed=None)))
model2dn.add(BatchNormalization())
model2dn.add(Dropout(0.5))

model2dn.add(Dense(output_dim, activation='softmax'))

# model Summary
print("model summary :- \n", model2dn.summary())

# Compiling the model
model2dn.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history2dn = model2dn.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type)                                | Output Shape | Param # |
|---|--------------|---------|
| dense_7 (Dense)                             | (None, 375)  | 294375  |
| batch_normalization_1 (Batch Normalization) | (None, 375)  | 1500    |
| dropout_1 (Dropout)                         | (None, 375)  | 0       |
| dense_8 (Dense)                             | (None, 112)  | 42112   |
| batch_normalization_2 (Batch Normalization) | (None, 112)  | 448     |

|                     |             |      |
|---------------------|-------------|------|
| dropout_2 (Dropout) | (None, 112) | 0    |
| dense_9 (Dense)     | (None, 10)  | 1130 |

---

Total params: 339,565  
 Trainable params: 338,591  
 Non-trainable params: 974

---

model summary :-

```

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 179us/step - loss: 0.4790 - acc: 0.8549 - val_loss: 0.1542 - val_acc: 0.9516
Epoch 2/20
60000/60000 [=====] - 9s 152us/step - loss: 0.2275 - acc: 0.9320 - val_loss: 0.1145 - val_acc: 0.9659
Epoch 3/20
60000/60000 [=====] - 9s 153us/step - loss: 0.1773 - acc: 0.9467 - val_loss: 0.0980 - val_acc: 0.9693
Epoch 4/20
60000/60000 [=====] - 10s 168us/step - loss: 0.1521 - acc: 0.9540 - val_loss: 0.0871 - val_acc: 0.9734
Epoch 5/20
60000/60000 [=====] - 9s 150us/step - loss: 0.1332 - acc: 0.9601 - val_loss: 0.0823 - val_acc: 0.9754
Epoch 6/20
60000/60000 [=====] - 9s 151us/step - loss: 0.1187 - acc: 0.9638 - val_loss: 0.0872 - val_acc: 0.9731
Epoch 7/20
60000/60000 [=====] - 9s 150us/step - loss: 0.1115 - acc: 0.9651 - val_loss: 0.0727 - val_acc: 0.9767
Epoch 8/20
60000/60000 [=====] - 10s 158us/step - loss: 0.1049 - acc: 0.9677 - val_loss: 0.0706 - val_acc: 0.9785
Epoch 9/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0966 - acc: 0.9700 - val_loss: 0.0667 - val_acc: 0.9791
Epoch 10/20
60000/60000 [=====] - 10s 160us/step - loss: 0.0910 - acc: 0.9715 - val_loss: 0.0655 - val_acc: 0.9790
Epoch 11/20
60000/60000 [=====] - 10s 163us/step - loss: 0.0853 - acc: 0.9740 - val_loss: 0.0680 - val_acc: 0.9796
Epoch 12/20
60000/60000 [=====] - 11s 175us/step - loss: 0.0820 - acc: 0.9738 - val_loss: 0.0644 - val_acc: 0.9804
Epoch 13/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0805 - acc: 0.9754 - val_loss: 0.0633 - val_acc: 0.9806
Epoch 14/20
60000/60000 [=====] - 9s 157us/step - loss: 0.0779 - acc: 0.9760 - val_loss: 0.0681 - val_acc: 0.9801
Epoch 15/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0755 - acc: 0.9765 - val_loss: 0.0674 - val_acc: 0.9806
Epoch 16/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0697 - acc: 0.9779 - val_loss: 0.0620 - val_acc: 0.9821
Epoch 17/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0688 - acc: 0.9788 - val_loss: 0.0615 - val_acc: 0.9820
Epoch 18/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0643 - acc: 0.9794 - val_loss: 0.0630 - val_acc: 0.9808
Epoch 19/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0655 - acc: 0.9791 - val_loss: 0.0651 - val_acc: 0.9800
Epoch 20/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0615 - acc: 0.9808 - val_loss: 0.0583 - val_acc: 0.9818
  
```

In [23]:

```
x = list(range(1,nb_epoch+1))
```




```
# getting Val loss
vy = history2dn.history['val_loss']
# getting Train loss
ty = history2dn.history['loss']

# function call
plt_dynamic_model(x, vy, ty)

# Evaluating the model
model_score = model2dn.evaluate(X_test, Y_test, verbose=0)
print('Test score:', model_score[0])
print('Test accuracy:', model_score[1])

# saving train and test accuracy of the model
model2dn_test_acc = model_score[1]
model2dn_train_acc = history2dn.history['acc']
```



Test score: 0.05833159902372281  
Test accuracy: 0.9818

## 2. Softmax classifier with 3-hidden layers Without dropout and Batch Normalization

In [21]:

```
model3 = Sequential()

model3.add(Dense(465, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model3.add(Dense(141, activation='relu', kernel_initializer=he_normal(seed=None)))
model3.add(Dense(65, activation='relu', kernel_initializer=he_normal(seed=None)))

model3.add(Dense(output_dim, activation='softmax'))

print(model3.summary())
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history3 = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type)              | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense_26 (Dense)          | (None, 465)  | 365025  |
| dense_27 (Dense)          | (None, 141)  | 65706   |
| dense_28 (Dense)          | (None, 65)   | 9230    |
| dense_29 (Dense)          | (None, 10)   | 660     |
| Total params: 440,621     |              |         |
| Trainable params: 440,621 |              |         |

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 177us/step - loss: 0.2408 - acc: 0.9295 - val\_loss: 0.1255 - val\_acc: 0.9601

Epoch 2/20

60000/60000 [=====] - 9s 154us/step - loss: 0.0889 - acc: 0.9729 - val\_loss: 0.0912 - val\_acc: 0.9725

Epoch 3/20

60000/60000 [=====] - 10s 168us/step - loss: 0.0570 - acc: 0.9820 - val\_loss: 0.0684 - val\_acc: 0.9787

Epoch 4/20

60000/60000 [=====] - 10s 167us/step - loss: 0.0401 - acc: 0.9867 - val\_loss: 0.0881 - val\_acc: 0.9742

Epoch 5/20

60000/60000 [=====] - 10s 171us/step - loss: 0.0287 - acc: 0.9905 - val\_loss: 0.0784 - val\_acc: 0.9781

Epoch 6/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0243 - acc: 0.9924 - val\_loss: 0.0783 - val\_acc: 0.9783

Epoch 7/20

60000/60000 [=====] - 10s 173us/step - loss: 0.0206 - acc: 0.9931 - val\_loss: 0.0766 - val\_acc: 0.9798

Epoch 8/20

60000/60000 [=====] - 9s 152us/step - loss: 0.0168 - acc: 0.9946 - val\_loss: 0.0844 - val\_acc: 0.9791

Epoch 9/20

60000/60000 [=====] - 10s 159us/step - loss: 0.0191 - acc: 0.9937 - val\_loss: 0.0804 - val\_acc: 0.9788

Epoch 10/20

60000/60000 [=====] - 10s 165us/step - loss: 0.0135 - acc: 0.9957 - val\_loss: 0.0807 - val\_acc: 0.9809

Epoch 11/20

60000/60000 [=====] - 10s 162us/step - loss: 0.0142 - acc: 0.9953 - val\_loss: 0.0707 - val\_acc: 0.9823

Epoch 12/20

60000/60000 [=====] - 11s 178us/step - loss: 0.0114 - acc: 0.9964 - val\_loss: 0.0849 - val\_acc: 0.9812

Epoch 13/20

60000/60000 [=====] - 9s 156us/step - loss: 0.0098 - acc: 0.9969 - val\_loss: 0.0859 - val\_acc: 0.9797

Epoch 14/20

60000/60000 [=====] - 9s 156us/step - loss: 0.0117 - acc: 0.9962 - val\_loss: 0.0986 - val\_acc: 0.9789

Epoch 15/20

60000/60000 [=====] - 10s 169us/step - loss: 0.0135 - acc: 0.9955 - val\_loss: 0.0783 - val\_acc: 0.9831

Epoch 16/20

60000/60000 [=====] - 11s 183us/step - loss: 0.0099 - acc: 0.9968 - val\_loss: 0.0885 - val\_acc: 0.9825

Epoch 17/20

60000/60000 [=====] - 11s 186us/step - loss: 0.0083 - acc: 0.9972 - val\_loss: 0.0940 - val\_acc: 0.9814s:

Epoch 18/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0088 - acc: 0.9976 - val\_loss: 0.0995 - val\_acc: 0.9811

Epoch 19/20

60000/60000 [=====] - 10s 171us/step - loss: 0.0095 - acc: 0.9972 - val\_loss: 0.1119 - val\_acc: 0.9760

Epoch 20/20

60000/60000 [=====] - 8s 137us/step - loss: 0.0086 - acc: 0.9973 - val\_loss: 0.0825 - val\_acc: 0.9829

In [24]:


```
x = list(range(1,nb_epoch+1))

# getting Val loss
vy = history3.history['val_loss']
# getting Train loss
ty = history3.history['loss']

# function call
plt_dynamic_model(x, vy, ty)
```

```
# Evaluating the model
model_score = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', model_score[0])
print('Test accuracy:', model_score[1])

# saving train and test accuracy of the model
model3_test_acc = model_score[1]
model3_train_acc = history3.history['acc']
```



Test score: 0.08253725015399864  
Test accuracy: 0.9829

## Softmax classifier with 3-hidden layers With Dropout and Batch Normalization

In [10]:

```
model3dn = Sequential()

model3dn.add(Dense(465, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model3dn.add(BatchNormalization())
model3dn.add(Dropout(0.5))

model3dn.add(Dense(141, activation='relu', kernel_initializer=he_normal(seed=None)))
model3dn.add(BatchNormalization())
model3dn.add(Dropout(0.5))

model3dn.add(Dense(65, activation='relu', kernel_initializer=he_normal(seed=None)))
model3dn.add(BatchNormalization())
model3dn.add(Dropout(0.5))

model3dn.add(Dense(output_dim, activation='softmax'))
print(model3dn.summary())
model3dn.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history3dn = model3dn.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type)                                | Output Shape | Param # |
|---|--------------|---------|
| dense_10 (Dense)                            | (None, 465)  | 365025  |
| batch_normalization_3 (Batch Normalization) | (None, 465)  | 1860    |
| dropout_3 (Dropout)                         | (None, 465)  | 0       |
| dense_11 (Dense)                            | (None, 141)  | 65706   |
| batch_normalization_4 (Batch Normalization) | (None, 141)  | 564     |
| dropout_4 (Dropout)                         | (None, 141)  | 0       |

|   |            |      |
|---|------------|------|
| dense_12 (Dense)                            | (None, 65) | 9230 |
| batch_normalization_5 (Batch Normalization) | (None, 65) | 260  |
| dropout_5 (Dropout)                         | (None, 65) | 0    |
| dense_13 (Dense)                            | (None, 10) | 660  |
| =====                                       |            |      |
| Total params: 443,305                       |            |      |
| Trainable params: 441,963                   |            |      |
| Non-trainable params: 1,342                 |            |      |

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 15s 249us/step - loss: 0.6878 - acc: 0.7869 - val\_loss: 0.1866 - val\_acc: 0.9435

Epoch 2/20

60000/60000 [=====] - 12s 198us/step - loss: 0.2941 - acc: 0.9160 - val\_loss: 0.1342 - val\_acc: 0.9599

Epoch 3/20

60000/60000 [=====] - 12s 196us/step - loss: 0.2237 - acc: 0.9382 - val\_loss: 0.1131 - val\_acc: 0.9666

Epoch 4/20

60000/60000 [=====] - 12s 194us/step - loss: 0.1874 - acc: 0.9471 - val\_loss: 0.0936 - val\_acc: 0.9719

Epoch 5/20

60000/60000 [=====] - 12s 198us/step - loss: 0.1680 - acc: 0.9526 - val\_loss: 0.0836 - val\_acc: 0.9739

Epoch 6/20

60000/60000 [=====] - 12s 195us/step - loss: 0.1513 - acc: 0.9578 - val\_loss: 0.0823 - val\_acc: 0.9744

Epoch 7/20

60000/60000 [=====] - 11s 186us/step - loss: 0.1370 - acc: 0.9615 - val\_loss: 0.0824 - val\_acc: 0.9765

Epoch 8/20

60000/60000 [=====] - 11s 186us/step - loss: 0.1331 - acc: 0.9627 - val\_loss: 0.0792 - val\_acc: 0.9777

Epoch 9/20

60000/60000 [=====] - 11s 189us/step - loss: 0.1241 - acc: 0.9657 - val\_loss: 0.0701 - val\_acc: 0.9792

Epoch 10/20

60000/60000 [=====] - 11s 188us/step - loss: 0.1115 - acc: 0.9685 - val\_loss: 0.0680 - val\_acc: 0.9814

Epoch 11/20

60000/60000 [=====] - 12s 202us/step - loss: 0.1044 - acc: 0.9698 - val\_loss: 0.0720 - val\_acc: 0.9793

Epoch 12/20

60000/60000 [=====] - 11s 191us/step - loss: 0.1035 - acc: 0.9713 - val\_loss: 0.0702 - val\_acc: 0.9798

Epoch 13/20

60000/60000 [=====] - 11s 190us/step - loss: 0.0996 - acc: 0.9722 - val\_loss: 0.0714 - val\_acc: 0.9802

Epoch 14/20

60000/60000 [=====] - 12s 198us/step - loss: 0.0957 - acc: 0.9717 - val\_loss: 0.0608 - val\_acc: 0.9821

Epoch 15/20

60000/60000 [=====] - 12s 193us/step - loss: 0.0905 - acc: 0.9746 - val\_loss: 0.0658 - val\_acc: 0.9806

Epoch 16/20

60000/60000 [=====] - 12s 197us/step - loss: 0.0862 - acc: 0.9759 - val\_loss: 0.0633 - val\_acc: 0.9825

Epoch 17/20

60000/60000 [=====] - 12s 201us/step - loss: 0.0857 - acc: 0.9749 - val\_loss: 0.0631 - val\_acc: 0.9815

Epoch 18/20

60000/60000 [=====] - 12s 197us/step - loss: 0.0810 - acc: 0.9763 - val\_loss: 0.0597 - val\_acc: 0.9830

Epoch 19/20

60000/60000 [=====] - 12s 194us/step - loss: 0.0776 - acc: 0.9778 - val\_loss: 0.0616 - val\_acc: 0.9829

Epoch 20/20

60000/60000 [=====] - 12s 204us/step - loss: 0.0734 - acc: 0.9786 - val\_loss: 0.0572 - val\_acc: 0.9834

In [19]:

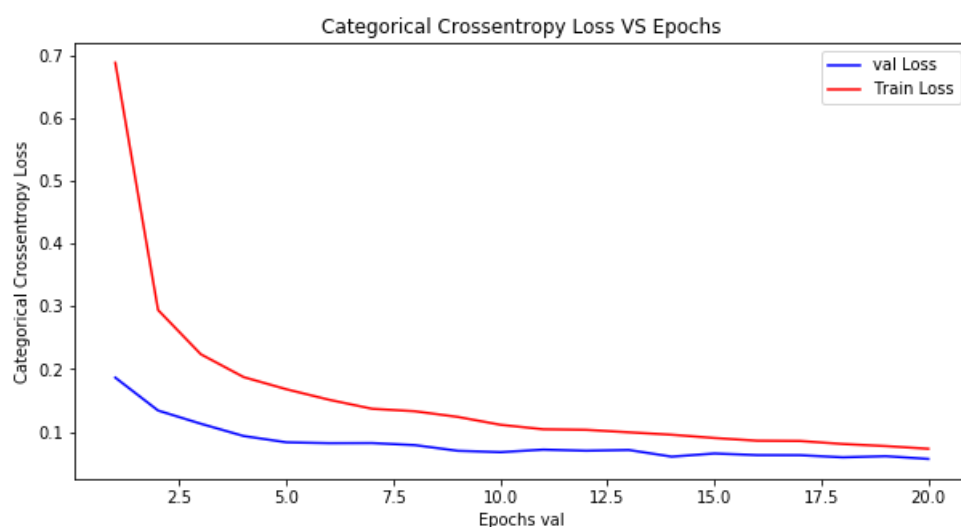
```
x = list(range(1,nb_epoch+1))

# getting Val loss
vy = history3dn.history['val_loss']
# getting Train loss
ty = history3dn.history['loss']

# function call
plt_dynamic_model(x, vy, ty)

# Evaluating the model
model_score = model3dn.evaluate(X_test, Y_test, verbose=0)
print('Test score:', model_score[0])
print('Test accuracy:', model_score[1])

# saving train and test accuracy of the model
model3dn_test_acc = model_score[1]
model3dn_train_acc = history3dn.history['acc']
```



Test score: 0.05717137544195284  
Test accuracy: 0.9834

### 3. Softmax classifier with 5-hidden layers without dropout and Batch Normalization

In [11]:

```
model5 = Sequential()

model5.add(Dense(513, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model5.add(Dense(252, activation='relu', kernel_initializer=he_normal(seed=None)))
model5.add(Dense(124, activation='relu', kernel_initializer=he_normal(seed=None)))
model5.add(Dense(68, activation='relu', kernel_initializer=he_normal(seed=None)))
model5.add(Dense(31, activation='relu', kernel_initializer=he_normal(seed=None)))

model5.add(Dense(output_dim, activation='softmax'))
print(model5.summary())
model5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history5 = model5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type)     | Output Shape | Param # |
|------------------|--------------|---------|
| dense_14 (Dense) | (None, 513)  | 402705  |
| dense_15 (Dense) | (None, 252)  | 129528  |
| dense_16 (Dense) | (None, 124)  | 31372   |

|                           |            |      |
|---------------------------|------------|------|
| dense_17 (Dense)          | (None, 68) | 8500 |
| dense_18 (Dense)          | (None, 31) | 2139 |
| dense_19 (Dense)          | (None, 10) | 320  |
| =====                     |            |      |
| Total params: 574,564     |            |      |
| Trainable params: 574,564 |            |      |
| Non-trainable params: 0   |            |      |

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 205us/step - loss: 0.2576 - acc: 0.9213 - val\_loss: 0.1284 - val\_acc: 0.9606s: 0.2580 - acc: 0.921

Epoch 2/20

60000/60000 [=====] - 12s 199us/step - loss: 0.0908 - acc: 0.9722 - val\_loss: 0.0920 - val\_acc: 0.9714

Epoch 3/20

60000/60000 [=====] - 11s 185us/step - loss: 0.0620 - acc: 0.9814 - val\_loss: 0.0863 - val\_acc: 0.9727

Epoch 4/20

60000/60000 [=====] - 12s 199us/step - loss: 0.0492 - acc: 0.9841 - val\_loss: 0.0750 - val\_acc: 0.9788

Epoch 5/20

60000/60000 [=====] - 12s 207us/step - loss: 0.0350 - acc: 0.9891 - val\_loss: 0.0834 - val\_acc: 0.9766

Epoch 6/20

60000/60000 [=====] - 11s 187us/step - loss: 0.0320 - acc: 0.9899 - val\_loss: 0.0706 - val\_acc: 0.9808

Epoch 7/20

60000/60000 [=====] - 12s 207us/step - loss: 0.0243 - acc: 0.9920 - val\_loss: 0.0812 - val\_acc: 0.9806

Epoch 8/20

60000/60000 [=====] - 12s 201us/step - loss: 0.0247 - acc: 0.9922 - val\_loss: 0.0835 - val\_acc: 0.9785

Epoch 9/20

60000/60000 [=====] - 11s 182us/step - loss: 0.0203 - acc: 0.9936 - val\_loss: 0.0936 - val\_acc: 0.9785

Epoch 10/20

60000/60000 [=====] - 11s 191us/step - loss: 0.0189 - acc: 0.9939 - val\_loss: 0.0854 - val\_acc: 0.9794

Epoch 11/20

60000/60000 [=====] - 13s 223us/step - loss: 0.0184 - acc: 0.9939 - val\_loss: 0.0839 - val\_acc: 0.9804

Epoch 12/20

60000/60000 [=====] - 13s 215us/step - loss: 0.0164 - acc: 0.9947 - val\_loss: 0.0763 - val\_acc: 0.9820

Epoch 13/20

60000/60000 [=====] - 13s 218us/step - loss: 0.0134 - acc: 0.9960 - val\_loss: 0.1041 - val\_acc: 0.9769

Epoch 14/20

60000/60000 [=====] - 11s 188us/step - loss: 0.0152 - acc: 0.9954 - val\_loss: 0.0660 - val\_acc: 0.9843

Epoch 15/20

60000/60000 [=====] - 11s 184us/step - loss: 0.0133 - acc: 0.9959 - val\_loss: 0.0851 - val\_acc: 0.9819

Epoch 16/20

60000/60000 [=====] - 11s 179us/step - loss: 0.0127 - acc: 0.9961 - val\_loss: 0.0964 - val\_acc: 0.9785 -

Epoch 17/20

60000/60000 [=====] - 11s 180us/step - loss: 0.0104 - acc: 0.9969 - val\_loss: 0.0784 - val\_acc: 0.9825

Epoch 18/20

60000/60000 [=====] - 11s 183us/step - loss: 0.0107 - acc: 0.9969 - val\_loss: 0.1170 - val\_acc: 0.9781

Epoch 19/20

60000/60000 [=====] - 12s 194us/step - loss: 0.0110 - acc: 0.9966 - val\_loss: 0.1014 - val\_acc: 0.9786

Epoch 20/20

60000/60000 [=====] - 11s 189us/step - loss: 0.0098 - acc: 0.9972 - val\_loss: 0.1006 - val\_acc: 0.9809

In [18]:

```

x = list(range(1,nb_epoch+1))


# getting Val loss
vy = history5.history['val_loss']
# getting Train loss
ty = history5.history['loss']

# function call
plt_dynamic_model(x, vy, ty)

# Evaluating the model
model_score = model5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', model_score[0])
print('Test accuracy:', model_score[1])

# saving train and test accuracy of the model
model5_test_acc = model_score[1]
model5_train_acc = history5.history['acc']

```



Test score: 0.10062611548545246  
Test accuracy: 0.9809

## Softmax classifier with 5-hidden layers With Dropout and Batch Normalisation

In [12]:

```

model5dn = Sequential()

model5dn.add(Dense(513, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model5dn.add(BatchNormalization())
model5dn.add(Dropout(0.5))

model5dn.add(Dense(252, activation='relu', kernel_initializer=he_normal(seed=None)))
model5dn.add(BatchNormalization())
model5dn.add(Dropout(0.5))

model5dn.add(Dense(124, activation='relu', kernel_initializer=he_normal(seed=None)))
model5dn.add(BatchNormalization())
model5dn.add(Dropout(0.5))

model5dn.add(Dense(68, activation='relu', kernel_initializer=he_normal(seed=None)))
model5dn.add(BatchNormalization())
model5dn.add(Dropout(0.5))

model5dn.add(Dense(31, activation='relu', kernel_initializer=he_normal(seed=None)))
model5dn.add(BatchNormalization())
model5dn.add(Dropout(0.5))

model5dn.add(Dense(output_dim, activation='softmax'))
print(model5dn.summary())
model5dn.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history5dn = model5dn.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, vali

```

```
validation_data=(x_test, y_test))
```

| Layer (type)                                 | Output Shape | Param # |
|--|--------------|---------|
| dense_20 (Dense)                             | (None, 513)  | 402705  |
| batch_normalization_6 (Batch Normalization)  | (None, 513)  | 2052    |
| dropout_6 (Dropout)                          | (None, 513)  | 0       |
| dense_21 (Dense)                             | (None, 252)  | 129528  |
| batch_normalization_7 (Batch Normalization)  | (None, 252)  | 1008    |
| dropout_7 (Dropout)                          | (None, 252)  | 0       |
| dense_22 (Dense)                             | (None, 124)  | 31372   |
| batch_normalization_8 (Batch Normalization)  | (None, 124)  | 496     |
| dropout_8 (Dropout)                          | (None, 124)  | 0       |
| dense_23 (Dense)                             | (None, 68)   | 8500    |
| batch_normalization_9 (Batch Normalization)  | (None, 68)   | 272     |
| dropout_9 (Dropout)                          | (None, 68)   | 0       |
| dense_24 (Dense)                             | (None, 31)   | 2139    |
| batch_normalization_10 (Batch Normalization) | (None, 31)   | 124     |
| dropout_10 (Dropout)                         | (None, 31)   | 0       |
| dense_25 (Dense)                             | (None, 10)   | 320     |
| Total params: 578,516                        |              |         |
| Trainable params: 576,540                    |              |         |
| Non-trainable params: 1,976                  |              |         |

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 19s 319us/step - loss: 1.4600 - acc: 0.5162 - val\_loss: 0.3729 - val\_acc: 0.9069

Epoch 2/20

60000/60000 [=====] - 18s 294us/step - loss: 0.6099 - acc: 0.8234 - val\_loss: 0.2034 - val\_acc: 0.9439

Epoch 3/20

60000/60000 [=====] - 16s 272us/step - loss: 0.4101 - acc: 0.8950 - val\_loss: 0.1691 - val\_acc: 0.9560

Epoch 4/20

60000/60000 [=====] - 17s 276us/step - loss: 0.3345 - acc: 0.9194 - val\_loss: 0.1409 - val\_acc: 0.9644

Epoch 5/20

60000/60000 [=====] - 16s 274us/step - loss: 0.2907 - acc: 0.9314 - val\_loss: 0.1330 - val\_acc: 0.9682

Epoch 6/20

60000/60000 [=====] - 16s 269us/step - loss: 0.2616 - acc: 0.9389 - val\_loss: 0.1170 - val\_acc: 0.9709

Epoch 7/20

60000/60000 [=====] - 16s 265us/step - loss: 0.2326 - acc: 0.9469 - val\_loss: 0.1168 - val\_acc: 0.9729

Epoch 8/20

60000/60000 [=====] - 16s 273us/step - loss: 0.2174 - acc: 0.9508 - val\_loss: 0.1052 - val\_acc: 0.9750

Epoch 9/20

60000/60000 [=====] - 16s 270us/step - loss: 0.2037 - acc: 0.9538 - val\_loss: 0.1045 - val\_acc: 0.9743

Epoch 10/20

60000/60000 [=====] - 16s 268us/step - loss: 0.1946 - acc: 0.9564 - val\_loss: 0.0987 - val\_acc: 0.9768

Epoch 11/20

60000/60000 [=====] - 16s 270us/step - loss: 0.1800 - acc: 0.9591 - val\_loss: 0.0994 - val\_acc: 0.9758

Epoch 12/20



```

Epoch 12/20
60000/60000 [=====] - 16s 270us/step - loss: 0.1758 - acc: 0.9611 - val_loss: 0.0921 - val_acc: 0.9783 - ETA: 2s - loss: 0 - ETA: 1s - loss: 0.1739 - acc: 0. - ETA: 1s - loss: 0.1
Epoch 13/20
60000/60000 [=====] - 16s 265us/step - loss: 0.1656 - acc: 0.9624 - val_loss: 0.0898 - val_acc: 0.9794
Epoch 14/20
60000/60000 [=====] - 16s 265us/step - loss: 0.1603 - acc: 0.9634 - val_loss: 0.0887 - val_acc: 0.9785
Epoch 15/20
60000/60000 [=====] - 17s 280us/step - loss: 0.1555 - acc: 0.9651 - val_loss: 0.0808 - val_acc: 0.9803
Epoch 16/20
60000/60000 [=====] - 17s 279us/step - loss: 0.1532 - acc: 0.9655 - val_loss: 0.0898 - val_acc: 0.9795
Epoch 17/20
60000/60000 [=====] - 16s 268us/step - loss: 0.1445 - acc: 0.9679 - val_loss: 0.0806 - val_acc: 0.9815
Epoch 18/20
60000/60000 [=====] - 16s 269us/step - loss: 0.1384 - acc: 0.9698 - val_loss: 0.0776 - val_acc: 0.9822
Epoch 19/20
60000/60000 [=====] - 16s 271us/step - loss: 0.1328 - acc: 0.9707 - val_loss: 0.0866 - val_acc: 0.9810
Epoch 20/20
60000/60000 [=====] - 16s 274us/step - loss: 0.1320 - acc: 0.9704 - val_loss: 0.0807 - val_acc: 0.9825

```

In [17]:

```

x = list(range(1,nb_epoch+1))


# getting Val loss
vy = history5dn.history['val_loss']
# getting Train loss
ty = history5dn.history['loss']

# function call
plt_dynamic_model(x, vy, ty)

# Evaluating the model
model_score = model5dn.evaluate(X_test, Y_test, verbose=0)
print('Test score:', model_score[0])
print('Test accuracy:', model_score[1])

# saving train and test accuracy of the model
model5dn_test_acc = model_score[1]
model5dn_train_acc = history5dn.history['acc']

```



Test score: 0.0807438613414066  
 Test accuracy: 0.9825

## Conclusion

In [30]:

```
from prettytable import PrettyTable

print('Performance Table')
x = PrettyTable()
x.field_names = ["Models", "Train", "Test"]

x.add_row(["2-Layer softmax without Dropout and BN", model2_train_acc[-1], model2_test_acc])
x.add_row(["2-Layer softmax with Dropout and BN ", model2dn_train_acc[-1], model2dn_test_acc])
x.add_row(["3-Layer softmax without Dropout and BN", model3_train_acc[-1], model3_test_acc])
x.add_row(["3-Layer softmax with Dropout and BN", model3dn_train_acc[-1], model3dn_test_acc])
x.add_row(["5-Layer softmax without Dropout and BN", model5_train_acc[-1], model5_test_acc])
x.add_row(["5-Layer softmax with Dropout and BN", model5dn_train_acc[-1], model5dn_test_acc])

print(x)
```

Performance Table

| Models                                 | Train              | Test   |
|--|--------------------|--------|
| 2-Layer softmax without Dropout and BN | 0.9967166666666667 | 0.982  |
| 2-Layer softmax with Dropout and BN    | 0.9808166666666667 | 0.9818 |
| 3-Layer softmax without Dropout and BN | 0.9972833333333333 | 0.9829 |
| 3-Layer softmax with Dropout and BN    | 0.9785833333651225 | 0.9834 |
| 5-Layer softmax without Dropout and BN | 0.9971666666666666 | 0.9809 |
| 5-Layer softmax with Dropout and BN    | 0.9703666666666667 | 0.9825 |